Mini-project 3
CMPSCI 670, Fall 2016, UMass Amherst
Due: November 4, 5:00 PM
Instructor: Subhransu Maji
TA: Tsung-Yu Lin

Guidelines

Submission. Submit a *single pdf document* via moodle that includes your solutions, figures and printouts of code. For readability you may attach the code printouts at the end of the solutions. You could have 24 hours late submission with a 50% mark down. Late submission beyond 24 hours will not be given *any* credits.

Plagiarism. We might reuse problem set questions from previous years, covered by papers and webpages, we expect the students not to copy, refer to, or look at the solutions in preparing their answers. We expect students to want to learn and not google for answers.

Collaboration. The homework must be done individually, except where otherwise noted in the assignments. 'Individually' means each student must hand in their own answers, and each student must write their own code in the programming part of the assignment. It is acceptable, however, for students to collaborate in figuring out answers and helping each other solve the problems. We will be assuming that you will be taking the responsibility to make sure you personally understand the solution to any work arising from such a collaboration.

Using other programming languages. All of the starter code is in Matlab which is what we expect you to use. You are free to use other languages such as Octave or Python with the caveat that we won't be able to answer or debug non Matlab questions.

1 Scale-space blob detection [45 points]

Figure 1: Scale-space blobs detection

In this homework you will implement a simple panoramic image stitching algorithm. The first part is to detect blobs as keypoints. The algorithm is outlined as follows:

1. Build a Laplacian scale space, starting with some initial scale and going for n iterations:

    (a) Filter image with scale-normalized Laplacian at current scale.

    (b) Save the square of Laplacian response for current level of scale space.

    (c) Increase scale by a factor k.

2. Perform non-maximum suppression in scale space.

3. Display resulting circles at their characteristic scales for points above a threshold.

Test images. There are four images in the folder `data/blobs` compressed in `p3_data.zip` to test your code. Fig. 1 provides the output of "butterfly" example as a reference to debug your code. Keep in mind that your output may look different depending on your threshold, range of scales, and other implementation details.

Running the code. Download `p3_code.zip` and `p3_data.zip` from moodle and start from the entry code `evalBlobsDetection.m`. It calls the dummy implementation of `detectBlobs.m` and draws the blob with `drawBlobs.m` as Fig. 2. You need to implement the detail in `detectBlobs.m` which output a $n \times 4$ matrix with blob in each row in (x, y, radius, score).


Figure 2: Output of running the default implementation of evalBlobsDetection

The following are the hints to implement the function:

- Don't forget to convert images to grayscale (`rgb2gray` command) and double (`im2double`).

- For creating the Laplacian filter, use the `fspecial` function (check the options). Pay careful attention to setting the right filter size.

- You have to choose the initial scale, the factor k by which the scale is multiplied each time, and the number of levels in the scale space. I typically set the initial scale to 2, and use 10 to 15 levels in the scale pyramid. The multiplication factor should depend on the largest scale at which you want regions to be detected.

- You may want to use a three-dimensional array to represent your scale space. It would be declared as:

  scaleSpace = zeros(h,w,n); % [h,w] - dimensions of image, n - number of levels in scale space

  Then *scaleSpace(:,:,i)* would give you the i'th level of the scale space. Alternatively, if you are storing different levels of the scale pyramid at different resolutions, you may want to use a cell array, where each "slot" can accommodate a different data type or a matrix of different dimensions. Here is how you would use it:

  scaleSpace = cell(n,1); %creates a cell array with n "slots"
  scaleSpace{i} = myMatrix; % store a matrix at level i

- A point is the local maximum if it is higher than all its neighbors (each point has 26 neighbors in 3D). You may find functions `nlfilter`, `colfilt` or `ordfilt2` useful for doing non-maximum supression.

- You also have to set a threshold on the squared Laplacian response above which to report region detections. You should play around with different values and choose one you like best.

- To display the detected regions as circles, you can use the `drawBlobs.m` function. You should display the highest scoring 1000 detected blobs for each image. If your code returns fewer than 1000 blobs then you may have to lower your threshold. Don't forget that there is a multiplication factor that relates the scale at which a region is detected to the radius of the circle that most closely approximates the region.

1.a What to submit

To get full credit of this part, you have to

- Include your implementation of `detecBlobs.m`.
- Include the results of given test images.

2 Image stitching [35 points]

In this part you will implement the RANSAC algorithm to stitch two images. The input to the algorithm are two images which are related by an unknown transformation. You will use your blobs detector implemented in the first part to extract keypoints and extract feature descriptors on them. Your goal is to estimate an affine transformation using feature matching and RANSAC to produce a combined image.

Recall that the overall steps in the alignment algorithm are:

1. Detect keypoints in each image (part 1).

2. Extract SIFT features at each detected keypoints.

3. Match features based on pairwise distance.

4. Use RANSAC to estimate the best affine transformation.

5. Stitch the two images using the estimated transformation.

The provided codebase contains an implementation of Step 2 and Step 5. You will implement the remaining steps. The entry code for this homework is in `evalStitching.m`. The code loads two images which contain overlapping areas. Figure 3 shows the "hill" example included in the homework. More test examples are given in the `data/stitching` folder compressed in `p3_data.zip`.

Figure 3: Input images to the image stitching algorithm.

Below is an outline of the steps you need to implement the algorithm.

1. **[0 points] Detect keypoints** The first step is to detect blobs in each image independently. You already implemented this in the first part. Remember to set the initial scale and threshold accordingly to get enough keypoints.

2. **[0 points] Feature extraction.** The next step is to extract feature descriptors on detected keypoints. The code to compute SIFT is included in the code base. The file `compute_sift.m` takes an image and (x, y, radius) of blobs as input and return a matrix of size $N \times 128$ with 128 dimensional SIFT descriptor in each row. The function uses different external libraries which are included in `mex` folder depending on the operating system. The first few lines in `evalStitching.m` sets up the environment variables in Matlab based on your operating system.

3. **[10 points] Computing matches.** The next step is to compute matches between the two sets of features. Implement the function `matches = computeMatches(f1,f2)` that returns the best match of each feature `f1` to `f2` using the smallest sum-of-squared-differences. The input are two matrices `f1=dxN` and `f2=dxM` and the output is a array of size `Nx1` where each entry `matches(i)` ∈

4

`[0,1,...,M]` is the closest feature in `f2` to the i^{th} feature `f1(:,i)`, with `matches(i) = 0` indicating no matching.

Considering only nearest neighbors returns many false matches. To refine the matches, compute the ratio of distance from the closest feature to the distance from the second closest feature and reject the matches with the ratio greater than a threshold. David Lowe's paper suggest a threshold 0.8. Set `matches(i)=0` if i^{th} feature `f1(:,i)` fails on the ratio test.

You can visualize the matches using the `showMatches(im1, im2, c1, c2, matches)` function provided in the codebase. Figure 4 shows the output of my implementation for reference.


Figure 4: Matches between features.

4. **[25 points] Estimating affine transformation using RANSAC.** The matching in the previous step is noisy and contains many outliers. Using RANSAC, estimate the affine translation that agrees with most matches (inliers). Implement the function `[inliers, transf] = ransac(matches, c1, c2)`. Here inliners is the indices of the inliner matches, transf is the estimated affine translation of 2×3 matrix $[m_1, m_2, t_1; m_3, m_4, t_2]$ between c1 and c2 (See the notation in Lecture slides). You need at least three pairs of points to estimate the this. See lecture slides for details.

You can visualize the inlier matches using the `showMatches` function. Figure 5 shows the output of my implementation for reference. The estimated affine transformation for this example is given as follows for your reference:

$$\begin{bmatrix} 1.006 & 0.041 & 142.02 \\ -0.047 & 1.005 & -17.94 \end{bmatrix}$$

Figure 5: Inliers obtained from RANSAC.

stitched image

Figure 6: Visualization of output.

5. **[0 points] Visualizing the stitched image.** Merge the images using `mergeImages(im1,im2, transf)` function provided in the codebase. The result is shown as Figure 6.

2.a What to submit

To get full credit for this part you have to

- Include your implementation of `computeMatches.m` and `ransac.m`.

- Visualize the intermediate results and the output of each of the provided test image pairs. You can do this by setting the `exampleIndex` variable in the `evalStitching.m` file.

- Report the estimated affine transformation for given test images.

3 Extensions for extra credit

- To detect blobs at multiple scales, there are two strategies to complete the task: (1) iteratively filter the image with kernel of increasing size. (2) iteratively downsample the image and filter it with the kernel of fixed size. You will have to upsample the result or do some interpolation in order to find maxima in scale space for the second approach. Which approach is more efficient? Run the both implementations and compare the performance between them.

- Implement the difference-of-Gaussian pyramid as mentioned in class and described in David Lowe's paper. Compare the results and the running time to the direct Laplacian implementation.

- Implement the affine adaptation step to turn circular blobs into ellipses as shown in the lecture (just one iteration is sufficient). The selection of the correct window function is essential here. You should use a Gaussian window that is a factor of 1.5 or 2 larger than the characteristic scale of the blob. Note that the lecture slides show how to find the relative shape of the second moment ellipse, but not the absolute scale (i.e., the axis lengths are defined up to some arbitrary constant multiplier). A good choice for the absolute scale is to set the sum of the major and minor axis half-lengths to the diameter of the corresponding Laplacian circle. To display the resulting ellipses, you should modify my circle-drawing function or look for a better function in the MATLAB documentation or on the Internet.

- The Laplacian has a strong response not only at blobs, but also along edges. However, recall from the class lecture that edge points are not "repeatable". So, implement an additional thresholding step that computes the Harris response at each detected Laplacian region and rejects the regions that have only one dominant gradient orientation (i.e., regions along edges). If you have implemented the affine adaptation step, these would be the regions whose characteristic ellipses are close to being degenerate (i.e., one of the eigenvalues is close to zero). Show both "before" and "after" detection results.

- There are few difficult test images where affine transformation will not produce perfect results. There are certain kinds of 2D transformation cannot be modeled as affine (See section 2.1 in Richard Szeliski's book). Try to align two images with homography transformation, which is more general than affine, to improve the results.

If you complete any work for extra credit, be sure to clearly mark that work in your report, explain it and include the code.